

---

---

# 2024 Coding Challenge: Automated Object Detection and Counting

— Taisha Joseph | [Github](#) | [LinkedIn](#) —  
June 24, 2024


---

---

# Objective

- Develop a Python program to count and detect objects from images using computer vision techniques.
- The model should be able to recognize and count the following objects (**person, car, bicycle**) in multiple contexts.

# Methodology



Data Pre-processing (  
Pycocotools, torch, JSON,  
requests etc.)



YOLOv5

# Data Preprocessing

- Download training and validation images with annotations from COCO
- Filter images and annotations for the 3 categories of interest using Python library pycocotools
- **Optional:** balance the number of images per classes by:
  - Limiting the number of images per categories based on the category with the minimum number of images
  - Limiting the number per categories based on a predefined number of images (i.e., 500) to reduce training time

```
def filter_coco_annotations(input_annotation_file, output_annotation_file, image_dir, output_image_dir, categories_to_keep):
    coco = COCO(input_annotation_file)
    category_ids = coco.getCatIds(catNms=categories_to_keep)
    image_ids = coco.getImgIds(catIds=category_ids)

    filtered_annotations = []
    filtered_images = []

    ensure_dir_exists(output_image_dir)
    ensure_dir_exists(os.path.dirname(output_annotation_file))

    for img_id in image_ids:
        img_info = coco.loadImgs(img_id)[0]
        ann_ids = coco.getAnnIds(imgIds=img_id, catIds=category_ids)
        anns = coco.loadAnns(ann_ids)

        filtered_images.append(img_info)
        filtered_annotations.extend(anns)

    # Copy image to output directory
    src_img_path = os.path.join(image_dir, img_info['file_name'])
    dst_img_path = os.path.join(output_image_dir, img_info['file_name'])
    shutil.copy(src_img_path, dst_img_path)

    filtered_data = {
        'images': filtered_images,
        'annotations': filtered_annotations,
        'categories': [cat for cat in coco.loadCats(category_ids)]
    }

    with open(output_annotation_file, 'w') as f:
        json.dump(filtered_data, f)
```

# Data Preprocessing

- Create indices for labels to be added to data configuration file
- Extract and normalize coordinates for bounding boxes from annotation file for each image according to YOLO format coordinates.
- Export YOLO-specific labels for each image

```
def convert_coco_to_yolo(annotations_file, labels_dir):
    # Load the COCO annotations file
    with open(annotations_file, 'r') as f:
        data = json.load(f)

    # Create a dictionary to map category IDs to category names
    categories = {cat['id']: cat['name'] for cat in data['categories']}

    # Create a dictionary to map category names to YOLO class indices
    category_to_index = {name: index for index, name in enumerate(categories.values())}

    # Ensure the labels directory exists
    os.makedirs(labels_dir, exist_ok=True)

    # Iterate over all annotations in the COCO dataset
    for ann in data['annotations']:
        image_id = ann['image_id']
        category_id = ann['category_id']
        bbox = ann['bbox']
        category_name = categories[category_id]

        # Skip categories that are not in the category_to_index dictionary
        if category_name not in category_to_index:
            continue

        # Get image information to calculate normalized bounding box coordinates
        image_info = next(img for img in data['images'] if img['id'] == image_id)
        image_width = image_info['width']
        image_height = image_info['height']

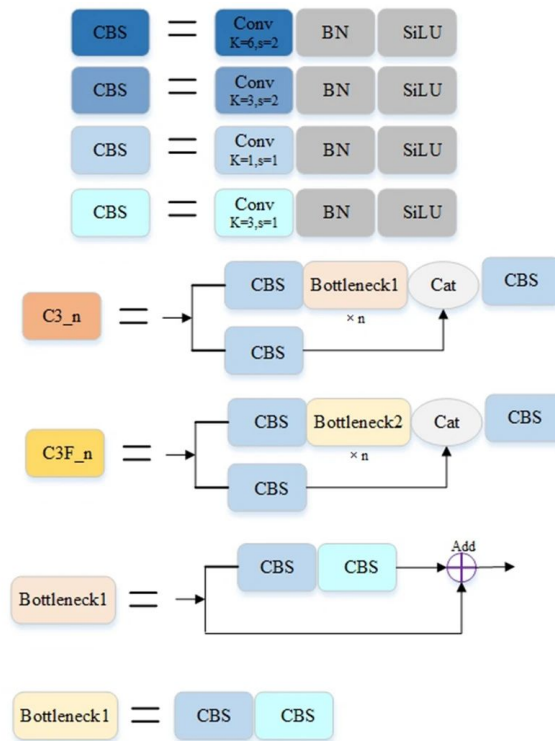
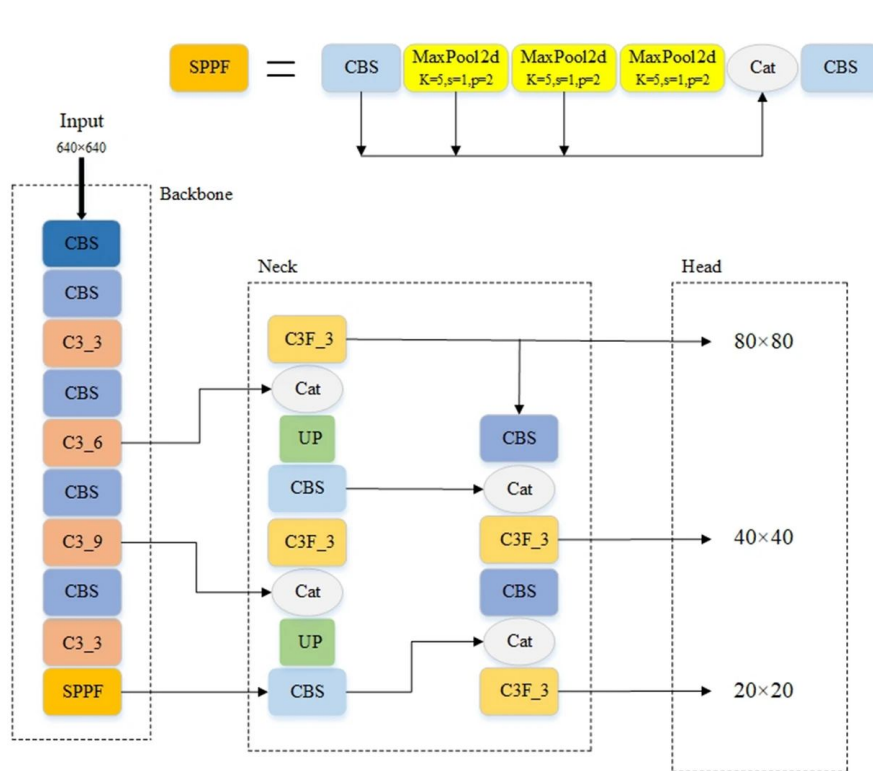
        # Calculate YOLO format coordinates (normalized)
        x_center = (bbox[0] + bbox[2] / 2) / image_width
        y_center = (bbox[1] + bbox[3] / 2) / image_height
        width = bbox[2] / image_width
        height = bbox[3] / image_height

        # Create the YOLO label string
        yolo_label = f"{category_to_index[category_name]} {x_center} {y_center} {width} {height}\n"

        # Determine the label file path based on the image file name (without extension)
        label_file_path = os.path.join(labels_dir, f"{image_info['file_name'].split('.')[0]}.txt")

        # Append the YOLO label to the label file
        with open(label_file_path, 'a') as label_file:
            label_file.write(yolo_label)
```

# YOLOv5 Network Architecture



# YOLOv5n Architecture

- Learning rate: 0.01
- Loss functions:
  - **Classification loss:**
    - Binary Cross-Entropy with Logits Loss (BCEWithLogitsLoss)
    - FocalLoss (variant of cross-entropy that handles class imbalance) if gamma (focusing param) > 0
  - **Bounding Box regression loss:**
    - Complete Intersection over Union (IoU) loss
- Activation function
  - **SiLU** (Swish) Activation Function (default)

```
# Parameters
nc: 3 # number of classes
depth_multiple: 0.33 # model depth multiple
width_multiple: 0.25 # layer channel multiple
anchors:
  - [10, 13, 16, 30, 33, 23] # P3/8
  - [30, 61, 62, 45, 59, 119] # P4/16
  - [116, 90, 156, 198, 373, 326] # P5/32

# YOLOv5 v6.0 backbone
backbone:
  # [from, number, module, args]
  [
    [-1, 1, Conv, [64, 6, 2, 2]], # 0-P1/2
    [-1, 1, Conv, [128, 3, 2]], # 1-P2/4
    [-1, 3, C3, [128]],
    [-1, 1, Conv, [256, 3, 2]], # 3-P3/8
    [-1, 6, C3, [256]],
    [-1, 1, Conv, [512, 3, 2]], # 5-P4/16
    [-1, 9, C3, [512]],
    [-1, 1, Conv, [1024, 3, 2]], # 7-P5/32
    [-1, 3, C3, [1024]],
    [-1, 1, SPPF, [1024, 5]], # 9
  ]

# YOLOv5 v6.0 head
head: [
  [-1, 1, Conv, [512, 1, 1]],
  [-1, 1, nn.Upsample, [None, 2, "nearest"]],
  [[-1, 6], 1, Concat, [1]], # cat backbone P4
  [-1, 3, C3, [512, False]], # 13

  [-1, 1, Conv, [256, 1, 1]],
  [-1, 1, nn.Upsample, [None, 2, "nearest"]],
  [[-1, 4], 1, Concat, [1]], # cat backbone P3
  [-1, 3, C3, [256, False]], # 17 (P3/8-small)

  [-1, 1, Conv, [256, 3, 2]],
  [[-1, 14], 1, Concat, [1]], # cat head P4
  [-1, 3, C3, [512, False]], # 20 (P4/16-medium)

  [-1, 1, Conv, [512, 3, 2]],
  [[-1, 10], 1, Concat, [1]], # cat head P5
  [-1, 3, C3, [1024, False]], # 23 (P5/32-large)

  [[17, 20, 23], 1, Detect, [nc, anchors]], # Detect(P3, P4, P5)
]
```

# Data Augmentations

- HSV augmentation  
(adjust the hue,  
saturation, and value  
(brightness) of images)
- Geometric  
transformations (random  
rotations, translations,  
scaling, shearing, etc.).
- Flipping and Mosaic  
(flipud, fliplr, mosaic)

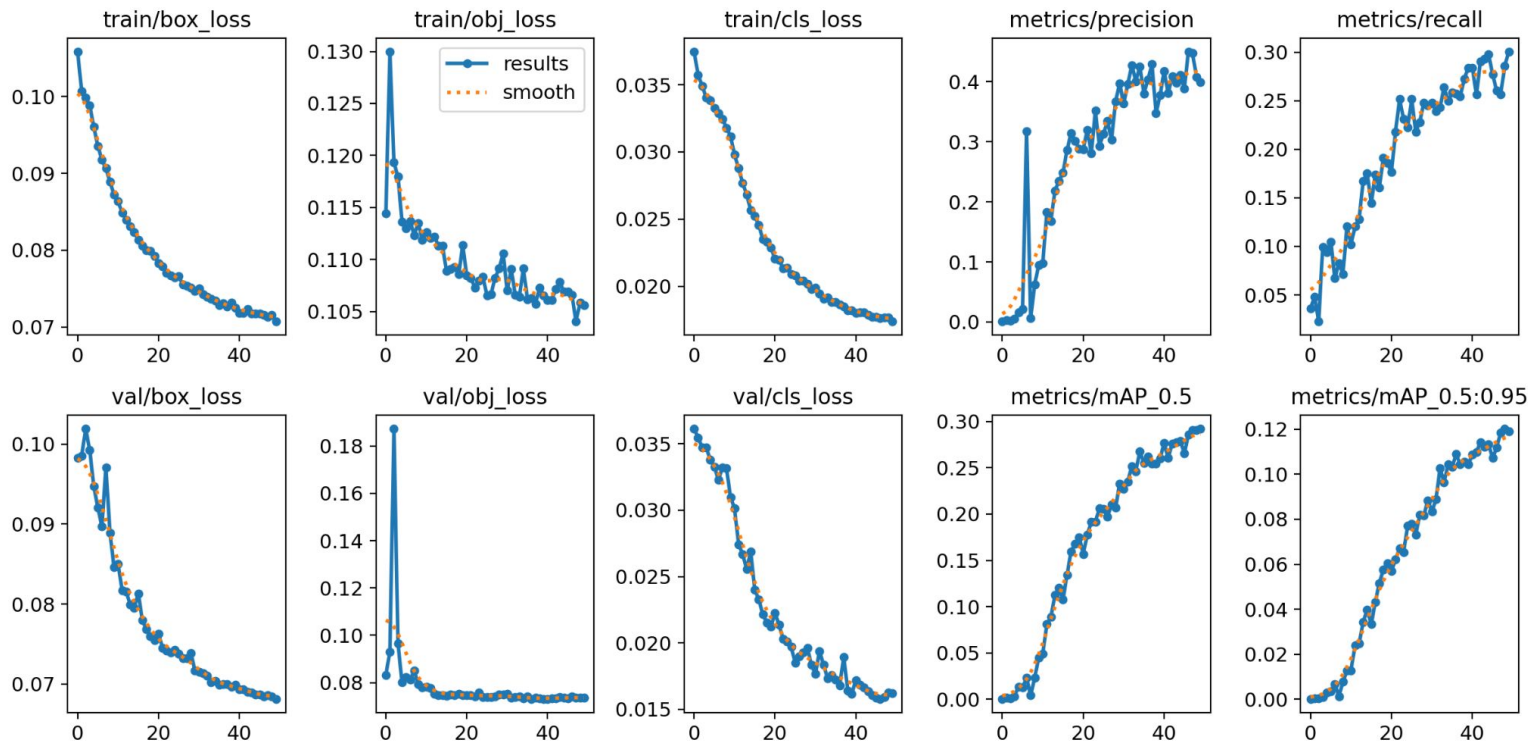
```
1 # Ultralytics YOLOv5 🚀, AGPL-3.0 license
2 # Hyperparameters for low-augmentation COCO training from scratch
3 # python train.py --batch 64 --cfg yolov5n6.yaml --weights '' --data coco.yaml --img 640 --epochs 300 --linear
4 # See tutorials for hyperparameter evolution https://github.com/ultralytics/yolov5#tutorials
5
6 lr0: 0.01 # initial learning rate (SGD=1E-2, Adam=1E-3)
7 lrf: 0.01 # final OneCycleLR learning rate (lr0 * lrf)
8 momentum: 0.937 # SGD momentum/Adam beta1
9 weight_decay: 0.0005 # optimizer weight decay 5e-4
10 warmup_epochs: 3.0 # warmup epochs (fractions ok)
11 warmup_momentum: 0.8 # warmup initial momentum
12 warmup_bias_lr: 0.1 # warmup initial bias lr
13 box: 0.05 # box loss gain
14 cls: 0.5 # cls loss gain
15 cls_pw: 1.0 # cls BCELoss positive_weight
16 obj: 1.0 # obj loss gain (scale with pixels)
17 obj_pw: 1.0 # obj BCELoss positive_weight
18 iou_t: 0.20 # IoU training threshold
19 anchor_t: 4.0 # anchor-multiple threshold
20 # anchors: 3 # anchors per output layer (0 to ignore)
21 fl_gamma: 0.0 # focal loss gamma (efficientDet default gamma=1.5)
22 hsv_h: 0.015 # image HSV-Hue augmentation (fraction)
23 hsv_s: 0.7 # image HSV-Saturation augmentation (fraction)
24 hsv_v: 0.4 # image HSV-Value augmentation (fraction)
25 degrees: 0.0 # image rotation (+/- deg)
26 translate: 0.1 # image translation (+/- fraction)
27 scale: 0.5 # image scale (+/- gain)
28 shear: 0.0 # image shear (+/- deg)
29 perspective: 0.0 # image perspective (+/- fraction), range 0-0.001
30 flipud: 0.0 # image flip up-down (probability)
31 fliplr: 0.5 # image flip left-right (probability)
32 mosaic: 1.0 # image mosaic (probability)
33 mixup: 0.0 # image mixup (probability)
34 copy_paste: 0.0 # segment copy-paste (probability)
```



# Training Object Detection Model

- 3 attempts with two models:
  1. Run 1:
    - a. Trained base YOLOv5n model from scratch for 50 epochs on full COCO 2017 training dataset
    - b. Validation using filtered COCO 2017 val dataset
  2. Run 2:
    - a. Added dropout regularization layers to improve performance of model on validation set and prevent overfitting.
    - b. Increased depth and width of network from model default to 0.5.
    - c. Trained modified model for 100 epochs on COCO 2017 training dataset limited by the size of the smallest category to balance the classes.
    - d. Validation using filtered COCO 2017 val dataset
  3. Run 3:
    - a. Same model architecture as run 2
    - b. Maintained balanced classes but limited the number of images in the training set to 500 for each class
    - c. Trained modified model for 200 epochs on limited COCO 2017 training dataset
    - d. Validation using filtered COCO 2017 val dataset

# Evaluation metrics



# Inference Results



## Inference

```
In [37]: !python detect.py --weights "runs/train/exp/weights/best.pt" --img 640 --conf 0.25 --source data/images
```

```
detect: weights=['runs/train/exp/weights/best.pt'], source=data/images, data=data/coco128.yaml, imgsz=[640, 640],  
conf_thres=0.25, iou_thres=0.45, max_det=1000, device=, view_img=False, save_txt=False, save_csv=False, save_conf=  
False, save_crop=False, nosave=False, classes=None, agnostic_nms=False, augment=False, visualize=False, update=False,  
project=runs/detect, name=exp, exist_ok=False, line_thickness=3, hide_labels=False, hide_conf=False, half=False,  
dnn=False, vid_stride=1
```

```
YOLOv5 v7.0-321-g3742ab49 Python-3.10.12 torch-2.3.0+cu121 CUDA:0 (NVIDIA L4, 22700MiB)
```

```
Fusing layers...
```

```
YOLOv5n_mod summary: 157 layers, 1763224 parameters, 0 gradients, 4.1 GFLOPs
```

```
WARNING ⚠ NMS time limit 0.550s exceeded
```

```
image 1/2 /content/yolov5/data/images/bus.jpg: 640x480 2 persons, 1 car, 1 bicycle, 98.8ms
```

```
image 2/2 /content/yolov5/data/images/zidane.jpg: 384x640 (no detections), 103.5ms
```

```
Speed: 0.5ms pre-process, 101.1ms inference, 282.3ms NMS per image at shape (1, 3, 640, 640)
```

```
Results saved to runs/detect/exp
```

# Inference Results



# How to Improve Performance?

- Increase number of epochs (from 50 to 100 or 200)
- Adding dropout layers
- Increase network depth and width to 0.5



# YOLOv5n' Architecture

- Learning rate: 0.01
- Loss functions:
  - **Classification loss:**
    - Binary Cross-Entropy with Logits Loss (BCEWithLogitsLoss)
    - FocalLoss (variant of cross-entropy that handles class imbalance) if gamma (focusing param) > 0
  - **Bounding Box regression loss:**
    - Complete Intersection over Union (IoU) loss
- Activation function
  - **SiLU** (Swish) Activation Function (default)
- Dropout regularization

```
# Parameters
nc: 3 # number of classes (modified)
depth_multiple: 0.5 # model depth multiple (modified)
width_multiple: 0.5 # layer channel multiple (modified)
anchors:
  - [10, 13, 16, 30, 33, 23] # P3/8
  - [30, 61, 62, 45, 59, 119] # P4/16
  - [116, 90, 156, 198, 373, 326] # P5/32

# YOLOv5 v6.0 backbone
backbone:
  # [from, number, module, args]
  [
    [-1, 1, Conv, [64, 6, 2, 2]], # 0-P1/2
    [-1, 1, Conv, [128, 3, 2]], # 1-P2/4
    [-1, 3, C3, [128]],
    [-1, 1, Conv, [256, 3, 2]], # 3-P3/8
    [-1, 6, C3, [256]],
    [-1, 1, Conv, [512, 3, 2]], # 5-P4/16
    [-1, 9, C3, [512]],
    [-1, 1, Conv, [1024, 3, 2]], # 7-P5/32
    [-1, 3, C3, [1024]],
    [-1, 1, nn.Dropout, [0.2]], # Adding Dropout with 20% probability
    [-1, 1, SPPF, [1024, 5]], # 9
  ]

# YOLOv5 v6.0 head
head: [
  [-1, 1, Conv, [512, 1, 1]],
  [-1, 1, nn.Upsample, [None, 2, "nearest"]],
  [[-1, 6], 1, Concat, [1]], # cat backbone P4
  [-1, 3, C3, [512, False]], # 13
  [-1, 1, nn.Dropout, [0.2]], # Adding Dropout with 20% probability

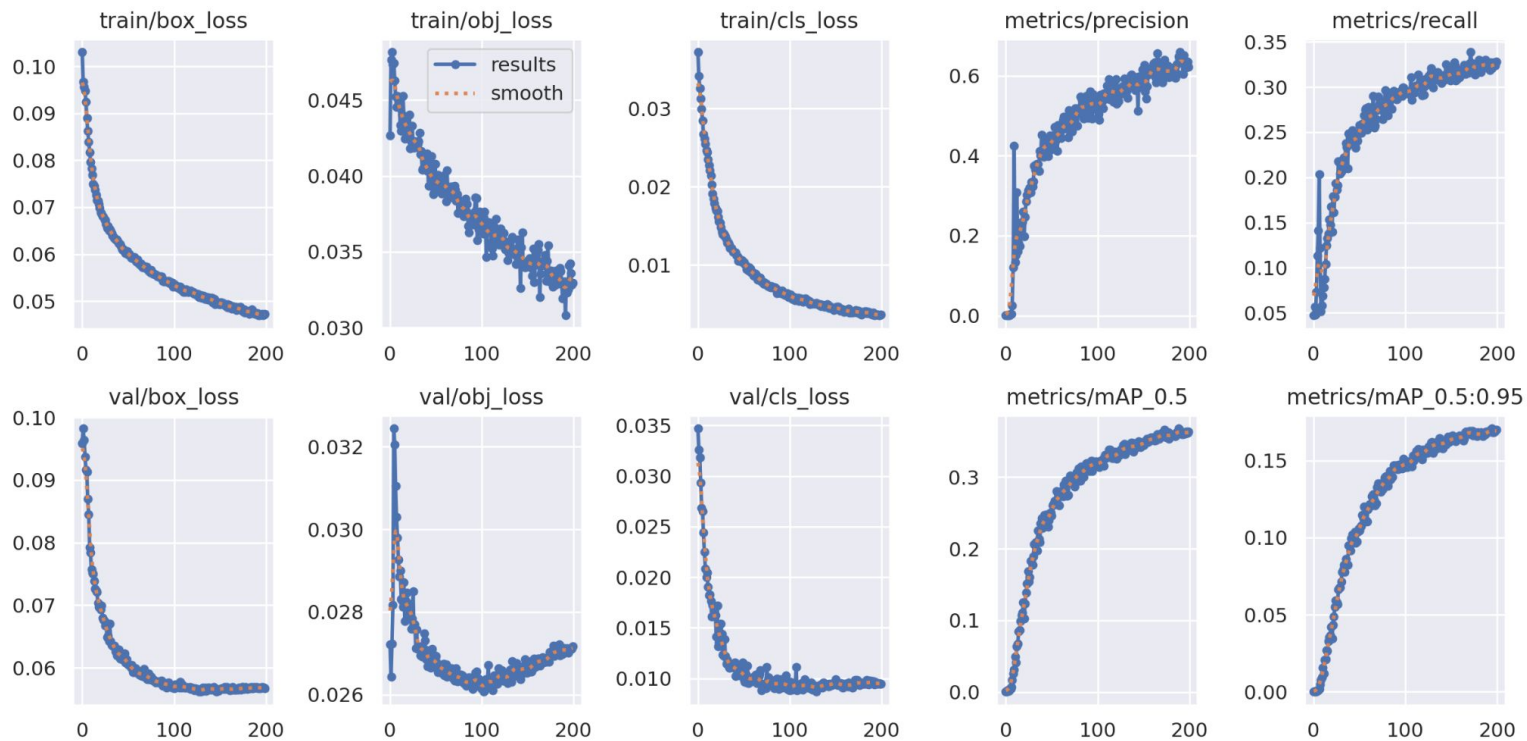
  [-1, 1, Conv, [256, 1, 1]],
  [-1, 1, nn.Upsample, [None, 2, "nearest"]],
  [[-1, 4], 1, Concat, [1]], # cat backbone P3
  [-1, 3, C3, [256, False]], # 17 (P3/8-small)
  [-1, 1, nn.Dropout, [0.2]], # Adding Dropout with 20% probability

  [-1, 1, Conv, [256, 3, 2]],
  [[-1, 14], 1, Concat, [1]], # cat head P4
  [-1, 3, C3, [512, False]], # 20 (P4/16-medium)
  [-1, 1, nn.Dropout, [0.3]], # Adding Dropout with 30% probability

  [-1, 1, Conv, [512, 3, 2]],
  [[-1, 10], 1, Concat, [1]], # cat head P5
  [-1, 3, C3, [1024, False]], # 23 (P5/32-large)
  [-1, 1, nn.Dropout, [0.5]], # Adding Dropout with 50% probability

  [[17, 20, 23], 1, Detect, [nc, anchors]], # Detect(P3, P4, P5)
]
```

# Evaluation metrics



# Inference Results

## Inference

```
In [33]: !python detect.py --weights {wts} --img 640 --conf 0.25 --source data/images
```

```
detect: weights=['runs/train/sknav2/exp4/weights/best.pt'], source=data/images, data=data/coco128.yaml, imgsz=[640, 640], conf_thres=0.25, iou_thres=0.45, max_det=1000, device=, view_img=False, save_txt=False, save_csv=False, save_conf=False, save_crop=False, nosave=False, classes=None, agnostic_nms=False, augment=False, visualize=False, update=False, project=runs/detect, name=exp, exist_ok=False, line_thickness=3, hide_labels=False, hide_conf=False, half=False, dnn=False, vid_stride=1
```

```
YOLOv5 🚀 v7.0-327-g098ce03f Python-3.10.12 torch-2.3.0+cu121 CUDA:0 (Tesla T4, 15102MiB)
```

```
Fusing layers...
```

```
YOLOv5n_mod summary: 202 layers, 9069352 parameters, 0 gradients, 20.2 GFLOPs
```

```
image 1/2 /content/yolov5/data/images/bus.jpg: 640x480 2 persons, 49.9ms
```

```
image 2/2 /content/yolov5/data/images/zidane.jpg: 384x640 2 persons, 50.9ms
```

```
Speed: 0.5ms pre-process, 50.4ms inference, 249.0ms NMS per image at shape (1, 3, 640, 640)
```

```
Results saved to runs/detect/exp3
```



# Inference Results



# Limitations/Potential Avenues for Improvements

- Train YOLOv5n' on full COCO 2017 train data instead of only a subset
- Increase network depth and width
- Increase number of epochs to 300
- Customize data augmentations prior to training by modifying appropriate hyp.yaml configuration file: i.e., increase image size, shearing, rotate, flip, zoom in/out etc.
- Decrease learning rate from 0.01 to 0.001

**Questions?**

# References/Resources

- <https://github.com/tavjo/computer-vision-project/tree/main>
- <http://cocodataset.org/>
- <https://github.com/cocodataset/cocoapi/tree/master>
- Jocher, G. (2020). YOLOv5 by Ultralytics (Version 7.0) [Computer software].  
<https://doi.org/10.5281/zenodo.3908559>
- Deng, L., Bi, L., Li, H. *et al.* Lightweight aerial image object detection algorithm based on improved YOLOv5s. *Sci Rep* 13, 7817 (2023).  
<https://doi.org/10.1038/s41598-023-34892-4>